

Krrish Kohli

5851 Beach Dr, Long Beach, CA, 90815 | [LinkedIn](#) | [GitHub](#) | krrishkohli15@gmail.com | (562) 564-2005

Summary

Honors Computer Science undergraduate (GPA 3.9, May 2027) building **production-minded software** across **backend APIs, data pipelines, and applied machine learning**. Experience delivering **user-facing prototypes**, designing **ETL workflows**, and applying **LLM/RAG + evaluation** to real datasets. Comfortable translating ambiguous stakeholder needs into scoped requirements, shipping working solutions, and documenting tradeoffs/results.

Education

California State University, Long Beach, CA, USA

Aug 2023 - May 2027

- Bachelor's degree in Computer Science (**Honors Program**)
- **GPA – 3.9 (Member of President's List)**
- **Coursework** - Digital Logic, Data Structures, Discrete Structures, System Programming, Algorithms, Computer Architecture, Software Engineering, Computer Security & Ethics, and Machine Learning & Artificial Intelligence.

Technical Skills

- **Languages:** Python, JavaScript, C++, C, TypeScript, C#, Java, SQL
- **Databases:** MongoDB, Neo4j, Relational Databases (SQL)
- **Web & App Frameworks:** React, Node.js, Express.js, EJS, HTML, CSS, REST APIs, Streamlit, Gradio
- **Data, AI & ML:** Pandas, NumPy, scikit-learn, SciPy, Matplotlib, Apache Spark, Apache Airflow, LLM APIs, RAG Workflows, Data Modeling, ETL concepts
- **Infrastructure & Tools:** Git, GitHub, Linux/Unix, Docker, CI/CD Basics, Cloud Fundamentals (AWS, Azure, GCP)

Work Experience

Data Engineering Intern | Springer Capital

Feb 2026 – May 2026

- Support design and maintenance of **data pipelines/ETL workflows**, including data cleaning, validation, and preprocessing for accuracy.
- Translate stakeholder data requirements into reliable datasets and **documented data flows** for repeatable downstream use.
- Monitor pipelines, troubleshoot failures, and improve storage/retrieval efficiency; research governance/security best practices.
- Use **SQL** for data manipulation and work with **Apache Spark** and **Apache Airflow** in the internship environment.

AI Studio Fellow | Brightside Health

Aug 2025 – Dec 2025

- Built and deployed a prototype enabling clinicians to explore evidence extracted from research papers (user-facing tool) using **Streamlit** and **Python**.
- Built a pipeline to convert papers into machine-readable text and extract entities (conditions/symptoms/treatments) into structured JSON.
- Implemented a **human-in-the-loop** + “**LLM-as-a-judge**” workflow to reduce hallucinations and improve fidelity to source material.
- Generated knowledge-graph-style representations from outputs to support navigation and decision workflows; documented usage and limitations.

Research Assistant | California State University, Long Beach

Nov 2024 – Dec 2025

- Developed predictive machine learning models to detect stroke survivor fall risks using motion data (500+ patients), improved prediction accuracy by 20%.
- Optimized feature selection to reduce compute time by 35% while maintaining 92% accuracy; evaluated sensitivity/error tradeoffs.
- Processed and analyzed gait/postural stability signals leveraging **scikit-learn**, **SciPy**, and **pandas** for feature extraction and analysis.

Projects

Clinical Interaction Knowledge Graph | [GitHub](#)

Aug 2025 - Dec 2025

- Built an end-to-end pipeline to ingest unstructured documents, extract structured entities/relationships, and persist results in Neo4j to support **AI for Knowledge Management**.
- Implemented a retrieval workflow to support **retrieval-augmented generation (RAG)** style analysis over domain content; emphasized deterministic outputs and schema validation.
- Created an evaluation harness to benchmark extraction/retrieval quality using metric-based checks; iterated based on measured failure modes (**automatic evaluation**, reliability discipline).
- **Languages and Frameworks:** Python, Neo4j, Docker, HTML, CSS, JavaScript, React.

FableFrog AI | [GitHub](#)

June 2024 - June 2024

- Built an AI storytelling app that generates narratives with an LLM and converts them to audio via ElevenLabs, delivered through a Gradio UI with end-to-end API integration.
- Improved reliability and latency by adding caching, rate limit aware retries with backoff and timeouts, instrumented structured logging to debug failures during demos.
- Refactored into modular components (generation, TTS, UI) with input validation and fallback behavior to handle edge cases and keep outputs consistent.
- **Languages and Frameworks:** Python, OpenAI API, ElevenLabs, OS, Gradio.